

Whole Genome Assembly and Alignment

Michael Schatz

Nov 6, 2012

SBU Graduate Genetics





Outline

1. Assembly theory
 1. Assembly by analogy
 2. De Bruijn and Overlap graph
 3. Coverage, read length, errors, and repeats
2. Genome assemblers
 1. Celera Assembler
3. Whole Genome Alignment with MUMmer
4. Review

Shredded Book Reconstruction

- Dickens accidentally shreds the first printing of A Tale of Two Cities
 - Text printed on 5 long spools

It was	the best	of times,	it was the worst	of times, it was the	age of wisdom, it was the	age of foolishness, ...
It was	the best	of times, it was the	the worst of times, it was the	the age of wisdom, it was the	the age of foolishness, ...	
It was	the best of times, it was	the worst of times, it	was the age of wisdom, i	it was the age of	foolishness, ...	
It was	the best of times, it was	the worst of times, it was the age of	wisdom, it was the age of	foolishness, ...		
It	was the best of times, it was the worst of	times, it was the age of	wisdom, it was the age of	foolishness, ...		

- How can he reconstruct the text?
 - 5 copies x 138,656 words / 5 words per fragment = 138k fragments
 - The short fragments from every copy are mixed together
 - Some fragments are identical

Greedy Reconstruction

It was the best of
age of wisdom, it was
best of times, it was
it was the age of
it was the age of
it was the worst of
of times, it was the
of times, it was the
of wisdom, it was the
the age of wisdom, it
the best of times, it
the worst of times, it
times, it was the age
times, it was the worst
was the age of wisdom,
was the age of foolishness,
was the best of times,
was the worst of times,
wisdom, it was the age
worst of times, it was

It was the best of
was the best of times,
the best of times, it
best of times, it was
of times, it was the
of times, it was the
times, it was the worst
times, it was the age

The repeated sequence make the correct reconstruction ambiguous

- It was the best of times, it was the [worst/age]

Model the assembly problem as a graph problem

de Bruijn Graph Construction

- $D_k = (V, E)$
 - V = All length- k subfragments ($k < l$)
 - E = Directed edges between consecutive subfragments
 - Nodes overlap by $k-1$ words

Original Fragment

It was the best of

Directed Edge

It was the best → was the best of

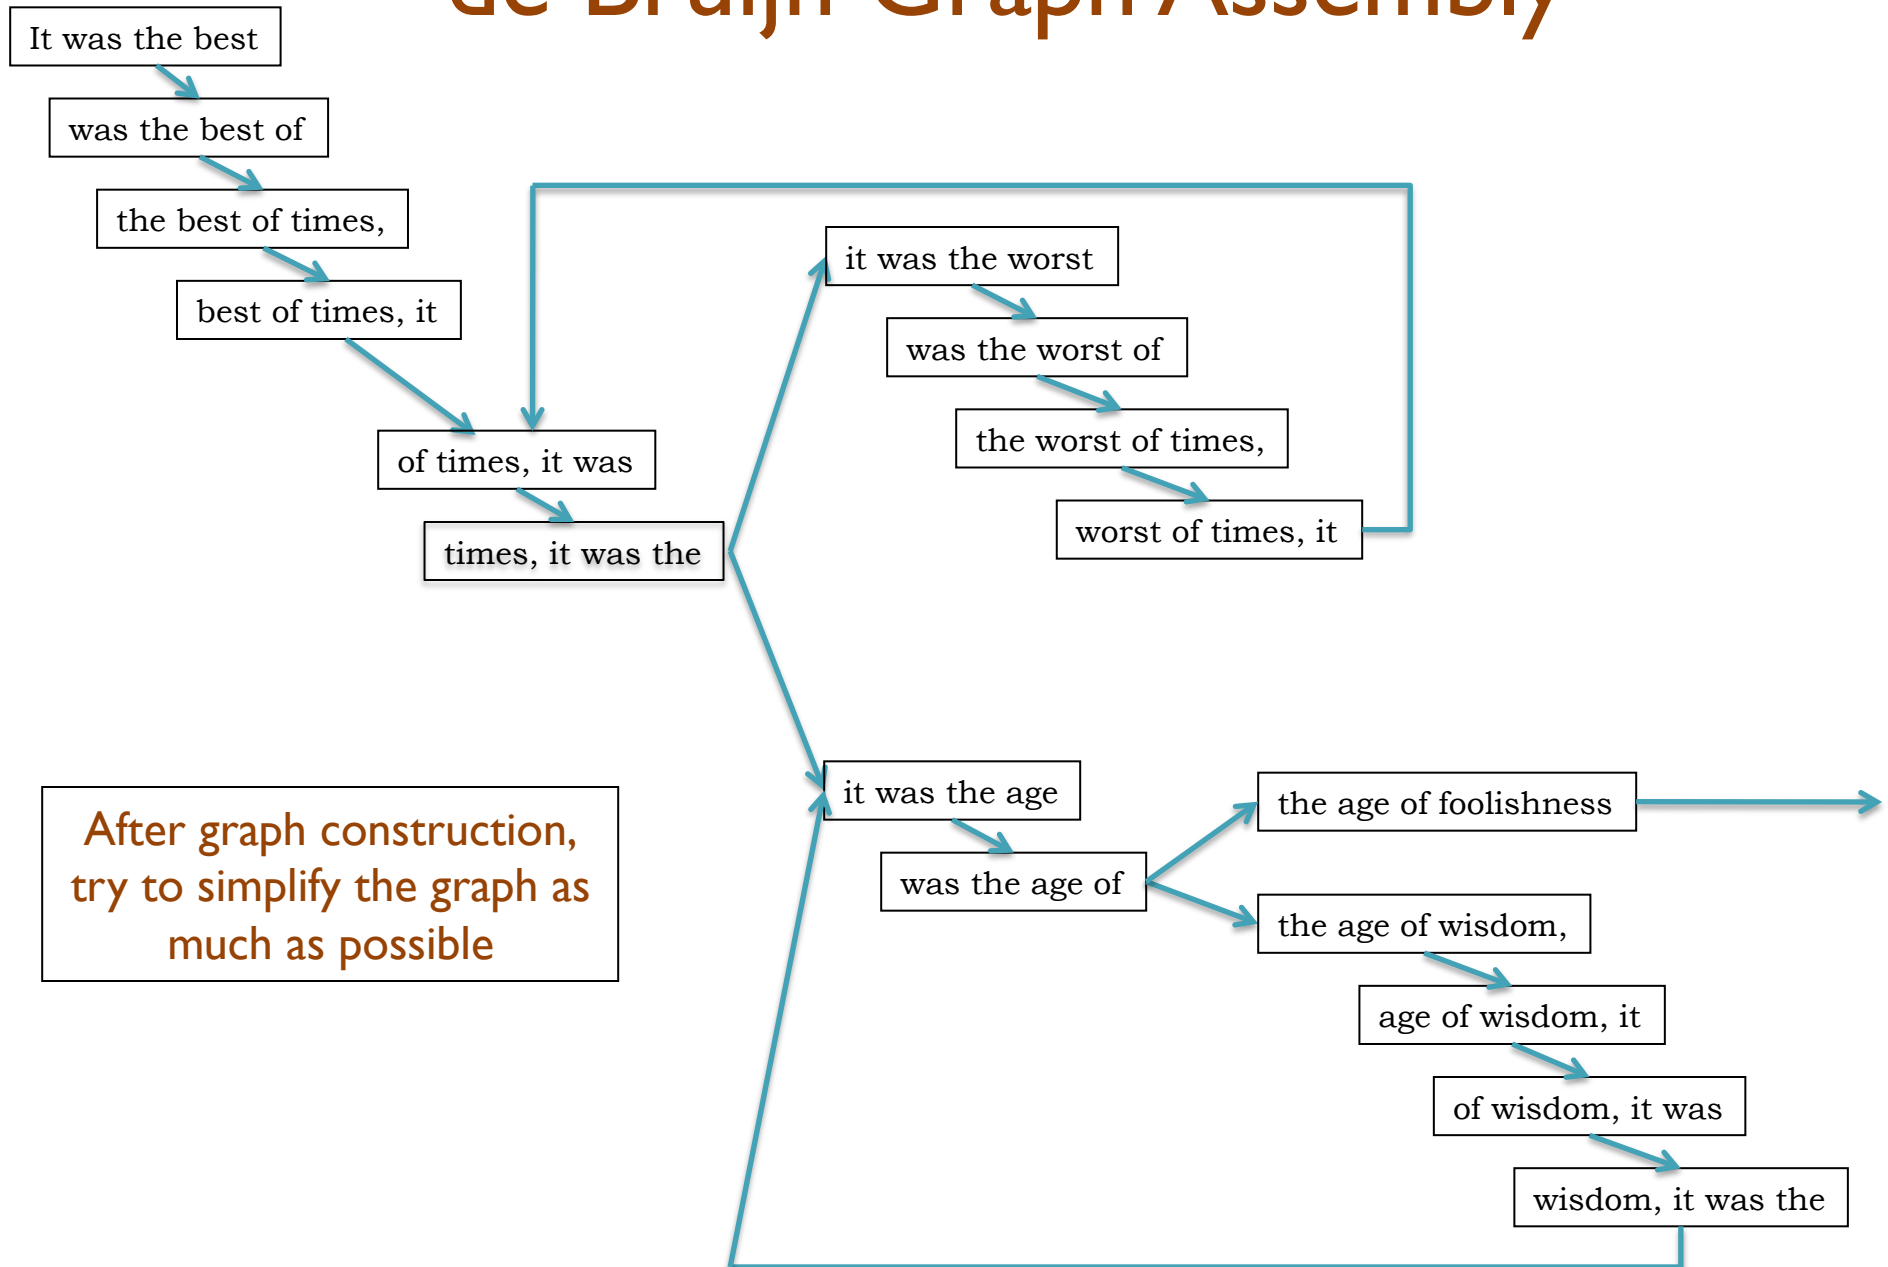
- Locally constructed graph reveals the global sequence structure
 - Overlaps between sequences implicitly computed

de Bruijn, 1946


Idury and Waterman, 1995

Pevzner, Tang, Waterman, 2001

de Bruijn Graph Assembly




de Bruijn Graph Assembly



Milestones in Genome Assembly



1977. Sanger *et al.*
1st Complete Organism
5375 bp



1995. Fleischmann *et al.*
1st Free Living Organism
TIGR Assembler. 1.8Mbp



1998. C.elegans SC
1st Multicellular Organism
BAC-by-BAC Phrap. 97Mbp



2000. Myers *et al.*
1st Large WGS Assembly.
Celera Assembler. 116 Mbp



2001. Venter *et al.*, IHGSC
Human Genome
Celera Assembler/GigaAssembler. 2.9 Gbp



2010. Li *et al.*
1st Large SGS Assembly.
SOAPdenovo 2.2 Gbp

Like Dickens, we must computationally reconstruct a genome from short fragments

Assembly Applications



- Novel genomes



- Metagenomes




- Sequencing assays
 - Structural variations
 - Transcript assembly
 - ...



Assembling a Genome


1. Shear & Sequence DNA




2. Construct assembly graph from overlapping reads

...AGCCTAGACCTACAGGATGCGCGACACGT
GGATGCGCGACACGTTCGCATATCCGGT...

3. Simplify assembly graph



4. Detangle graph with long reads, mates, and other links



Why are genomes hard to assemble?

1. Biological:

- (Very) High ploidy, heterozygosity, repeat content

2. Sequencing:

- (Very) large genomes, imperfect sequencing

3. Computational:

- (Very) Large genomes, complex structure


4. Accuracy:

- (Very) Hard to assess correctness



Ingredients for a good assembly


Coverage



High coverage is required

- Oversample the genome to ensure every base is sequenced with long overlaps between reads
- Biased coverage will also fragment assembly


Read Length



Reads & mates must be longer than the repeats

- Short reads will have **false overlaps** forming hairball assembly graphs
- With long enough reads, assemble entire chromosomes into contigs

Quality




Errors obscure overlaps

- Reads are assembled by finding kmers shared in pair of reads
- High error rate requires very short seeds, increasing complexity and forming assembly hairballs


Current challenges in *de novo* plant genome sequencing and assembly

Schatz MC, Witkowski, McCombie, VWR (2012) *Genome Biology*. 12:243


Illumina Sequencing by Synthesis




1. Prepare




2. Attach



3. Amplify



4. Image



5. Basecall

Metzker (2010) Nature Reviews Genetics 11:31-46

http://www.illumina.com/documents/products/techspotlights/techspotlight_sequencing.pdf

Paired-end and Mate-pairs

Paired-end sequencing


- Read one end of the molecule, flip, and read the other end
- Generate pair of reads separated by up to 500bp with inward orientation



Mate-pair sequencing

- Circularize long molecules (1-10kbp), shear into fragments, & sequence
- Mate failures create short paired-end reads

10kbp



2x100 @ ~10kbp (outies)




2x100 @ 300bp (innies)





Coverage

Typical contig coverage





Imagine raindrops on a sidewalk



Ix Sequencing





2x Sequencing




3x Sequencing




4x Sequencing




5x Sequencing




6x Sequencing



7x Sequencing




8x Sequencing



Coverage and Read Length

Idealized Lander-Waterman model

- Reads start at perfectly random positions
- Contig length is a function of coverage and read length
 - Short reads require much higher coverage to reach same expected contig length
- Need even high coverage for higher ploidy, sequencing errors, sequencing biases
 - Recommend 100x coverage




Assembly of Large Genomes using Second Generation Sequencing

Schatz MC, Delcher AL, Salzberg SL (2010) *Genome Research*. 20:1165-1173.

Two Paradigms for Assembly


de Bruijn Graph



Short read assemblers

- Repeats depends on word length
- Read coherency, placements lost
- Robust to high coverage

Overlap Graph



Long read assemblers



- Repeats depends on read length
- Read coherency, placements kept
- Tangled by high coverage

Assembly of Large Genomes using Second Generation Sequencing

Schatz MC, Delcher AL, Salzberg SL (2010) *Genome Research*. 20:1165-1173.

Unitigging / Unipathing

- After simplification and correction, compress graph down to its non-branching initial contigs
 - Aka “unitigs”, “unipaths”
 - Unitigs end because of (1) lack of coverage, (2) errors, (3) repeats, and (4) heterozygosity




Errors in the graph



(Chaisson, 2009)

Clip Tips	Pop Bubbles
<div>was the worst of times,</div> <div>was the worst of tymes,</div> <div>the worst of times, it</div>	<div>was the worst of times,</div> <div>was the worst of tymes,</div> <div>times, it was the age</div> <div>tymes, it was the age</div>
<div>the worst of tymes,</div> <div>was the worst of</div> <div>the worst of times,</div> <div>worst of times, it</div>	<div>tymes,</div> <div>was the worst of</div> <div>it was the age</div> <div>times,</div>

Repeats and Read Length



- Explore the relationship between read length and contig N50 size
 - Idealized assembly of read lengths: 25, 35, 50, 100, 250, 500, 1000
 - Contig/Read length relationship depends on specific repeat composition

Assembly Complexity of Prokaryotic Genomes using Short Reads.


Kingsford C, Schatz MC, Pop M (2010) *BMC Bioinformatics*. 11:21.

Repetitive regions

Repeat Type	Definition / Example	Prevalence
Low-complexity DNA / Microsatellites	$(b_1b_2\dots b_k)^N$ where $1 \leq k \leq 6$ CACACACACACACACACA	2%
SINEs (Short Interspersed Nuclear Elements)	<i>Alu</i> sequence (~280 bp) Mariner elements (~80 bp)	13%
LINEs (Long Interspersed Nuclear Elements)	~500 – 5,000 bp	21%
LTR (long terminal repeat) retrotransposons	Ty1-copia, Ty3-gypsy, Pao-BEL (~100 – 5,000 bp)	8%
Other DNA transposons		3%
Gene families & segmental duplications		4%

- Over 50% of mammalian genomes are repetitive
 - Large plant genomes tend to be even worse
 - Wheat: 16 Gbp; Pine: 24 Gbp

Repeats and Coverage Statistics




- If n reads are a uniform random sample of the genome of length G , we expect $k = n \Delta / G$ reads to start in a region of length Δ .
 - If we see many more reads than k (if the arrival rate is $> \lambda$), it is likely to be a collapsed repeat
 - Requires an accurate genome size estimate

$$\Pr(X = \text{copy}) = \binom{n}{k} \left(\frac{\Delta}{G} \right)^k \left(\frac{G - \Delta}{G} \right)^{n-k}$$

$$A(\Delta, k) = \ln \left(\frac{\Pr(1 - \text{copy})}{\Pr(2 - \text{copy})} \right) = \ln \left(\frac{\frac{(\Delta n / G)^k e^{-\Delta n / G}}{k!}}{\frac{(2\Delta n / G)^k e^{-2\Delta n / G}}{k!}} \right) = \frac{n\Delta}{G} - k \ln 2$$

Scaffolding

- Initial contigs (*aka* unipaths, unitigs) terminate at
 - *Coverage gaps*: especially extreme GC regions
 - *Conflicts*: sequencing errors, repeat boundaries
- Iteratively resolve longest, ‘most unique’ contigs
 - Both overlap graph and de Bruijn assemblers initially collapse repeats into single copies
 - Uniqueness measured by a statistical test on coverage



N50 size

Def: 50% of the genome is in contigs larger than N50

Example: 1 Mbp genome

50%



N50 size = 30 kbp

(300k+100k+45k+45k+30k = 520k >= 500kbp)

Note:

N50 values are only meaningful to compare when base genome size is the same in all cases

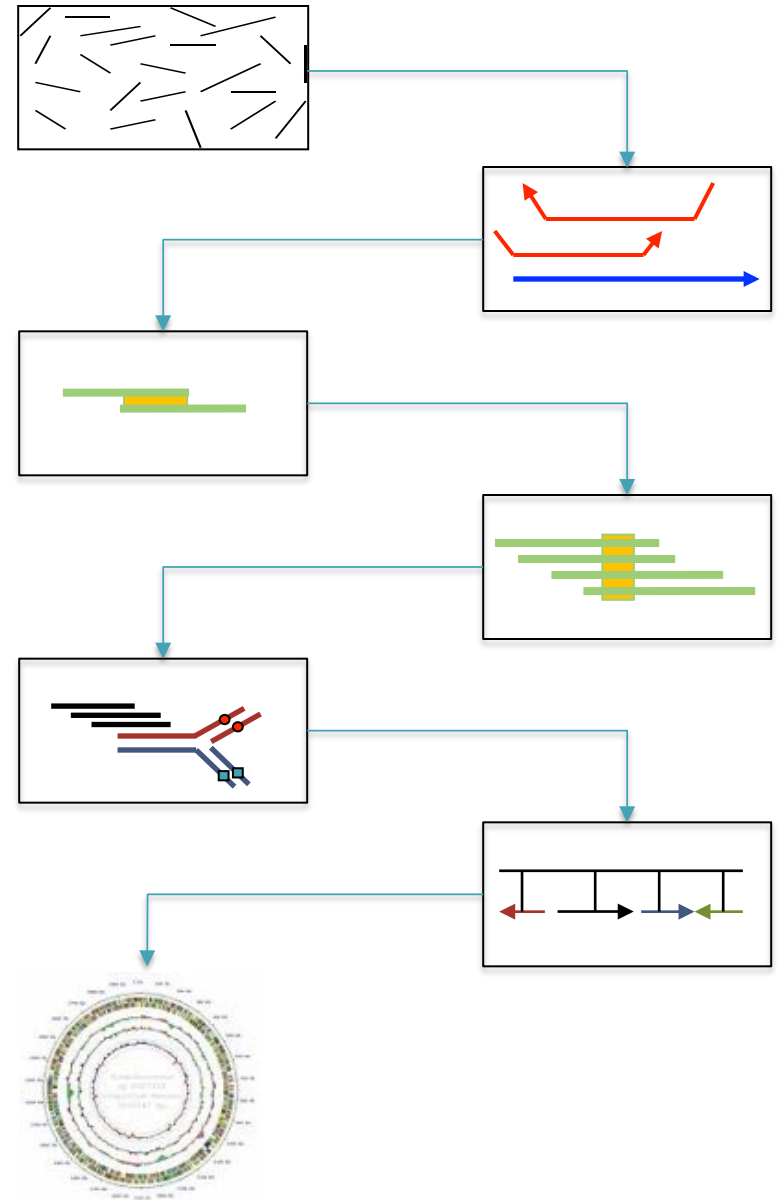


Genome assembly with the Celera Assembler

Celera Assembler

<http://wgs-assembler.sf.net>

1. Pre-overlap
 - Consistency checks
2. Trimming
 - Quality trimming & partial overlaps
3. Compute Overlaps
 - Find high quality overlaps
4. Error Correction
 - Evaluate difference in context of overlapping reads
5. Unitigging
 - Merge consistent reads
6. Scaffolding
 - Bundle mates, Order & Orient
7. Finalize Data
 - Build final consensus sequences



Hybrid Sequencing



Illumina

Sequencing by Synthesis

High throughput (60Gbp/day)
High accuracy (~99%)
Short reads (~100bp)



Pacific Biosciences


SMRT Sequencing

Lower throughput (600Mbp/day)
Lower accuracy (~85%)
Long reads (2-5kbp+)

SMRT Sequencing Data

Yeast
(Pre-release Chemistry / 2010)

65 SMRT cells
734,151 reads after filtering
Mean: 642.3 +/- 587.3
Median: 553 Max: 8,495



TTGTAAGCAGTTGAAAACATATGTGTGGATTTAGATAAAGAACATGAAAG
TTGTAAGCAGTTGAAAACATATGTGT-GATTTAG-ATAAAGAACATGGAAG

ATTATAAA-CAGTTGATCCATT-AGAAGA-AAACGCAAAGGCAGCTAGG
A-TATAAATCAGTTGATCCATTAGAA-AGAAACGC-AAAGGC-GCTAGG

CAACCTTGAAATGTAATCGCACTTGAAGAACAAGATTTTATTCCGCGCCCG
C-ACCTTG-ATGT-AT--CACTTGAAGAACAAGATTTTATTCCGCGCCCG

TACGAATCAAGATTCTGAAAACACAT-ATAACAACCTCCAAAA-CACAA
T-ACGAATC-AGATTCTGAAAACA-ATGAT----ACCTCCAAAAGCACAA

-AGGAGGGGAAAAGGGGGAATATCT-ATAAAGATTACAAATTAGA-TGA
GAGGAGG---AA-----GAATATCTGAT-AAAGATTACAAATT-GAGTGA

ACT-AATTCACAAATA-AATAACACTTTTA-ACAGAATTGAT-GGAA-GTT
ACTAAATTCACAA-ATAATAACACTTTTAGACAATAATTGATGGGAAGGTT

TCGGAGAGATCCAAAACAATGGGC-ATCGCCTTTGA-GTTAC-AATCAAA
TC-GAGAGATCC-AAACAAT-GGCGATCG-CTTTGACGTTACAAATCAAA

ATCCAGTGAAAAATATAATTTATGCAATCCAGGAACCTTATTCACAATTAG
ATCCAGT-GAAAAATATA--TTATGC-ATCCA-GAAGCTTATTCACAATTAG

Sample of 100k reads aligned with BLASR requiring >100bp alignment
Average overall accuracy: 83.7%, 11.5% insertions, 3.4% deletions, 1.4% mismatch

PacBio Error Correction


<http://wgs-assembler.sf.net>



I. Correction Pipeline

1. Map short reads to long reads
2. Trim long reads at coverage gaps
3. Compute consensus for each long read

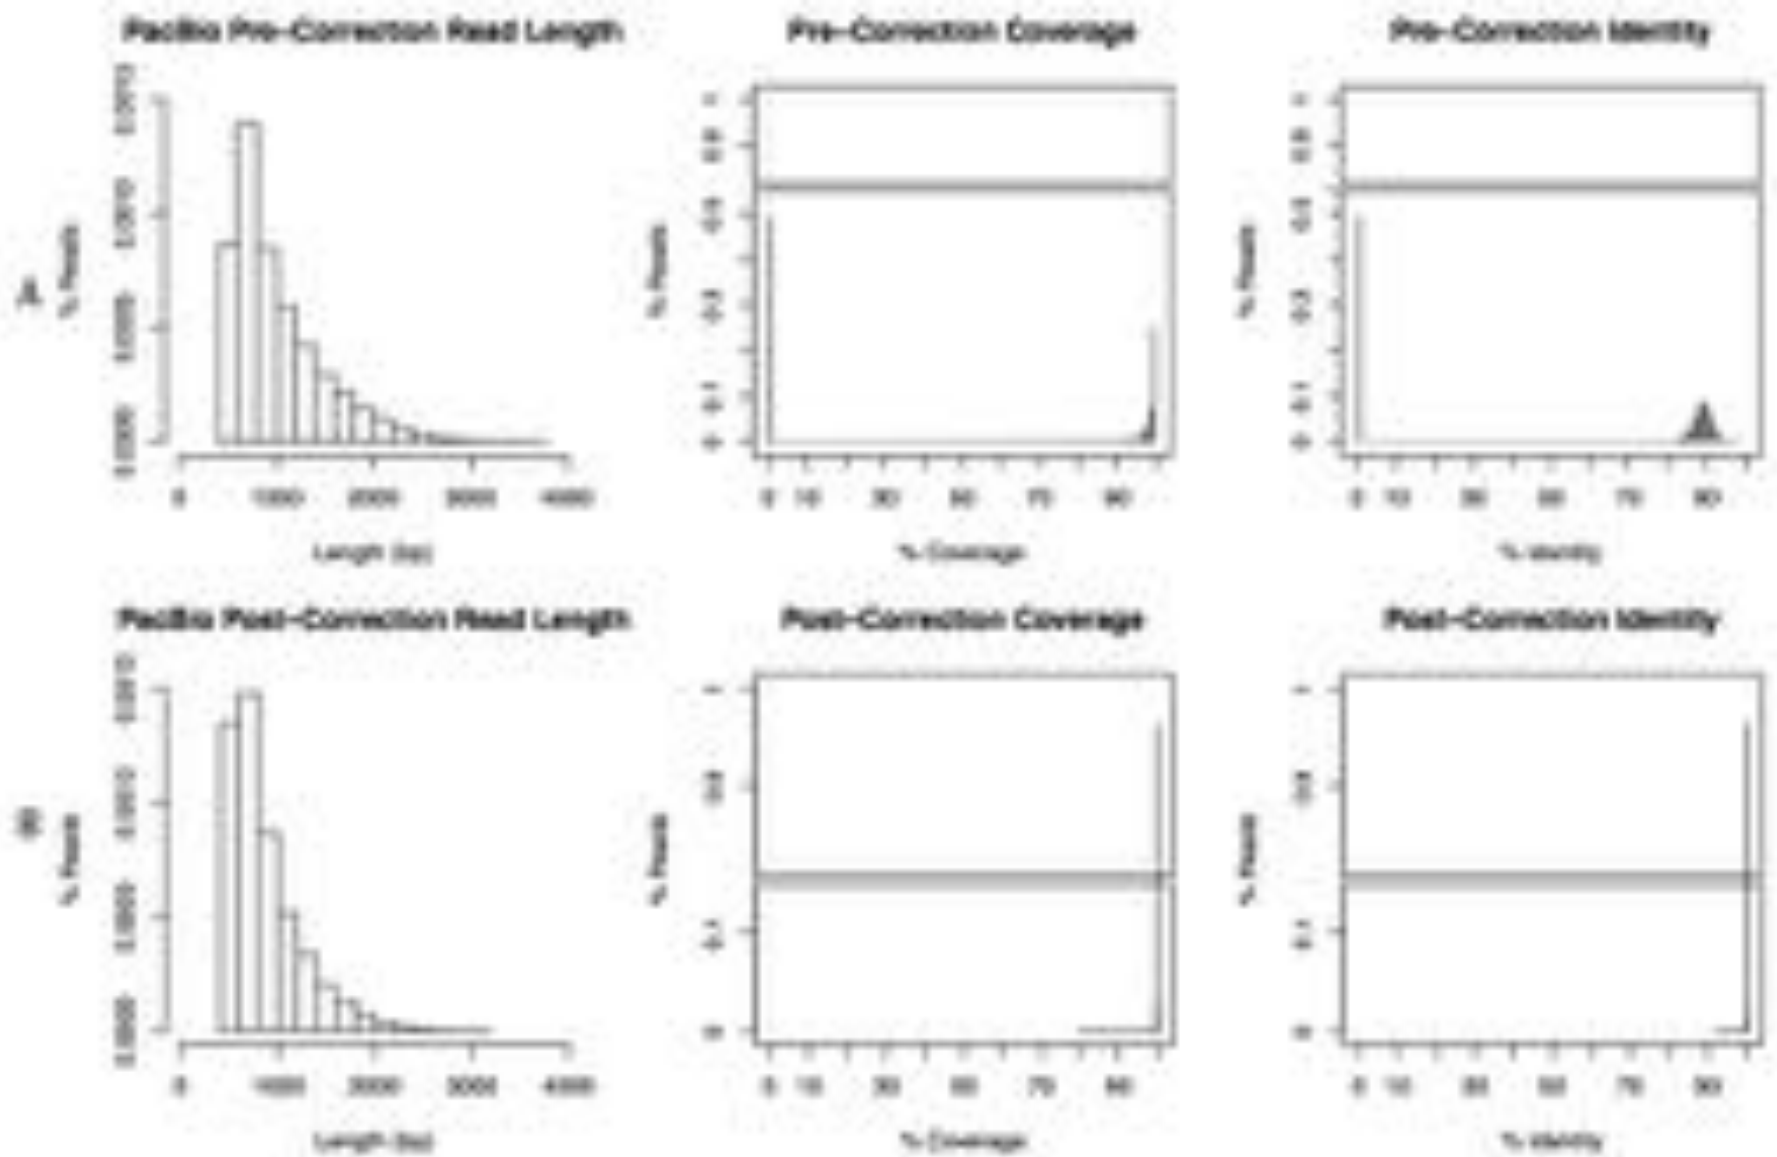
2. Error corrected reads can be easily assembled, aligned



Hybrid error correction and de novo assembly of single-molecule sequencing reads.


Koren, S, Schatz, MC, et al. (2012) *Nature Biotechnology*. doi:10.1038/nbt.2280

Error Correction Results



Correction results of 20x PacBio coverage of E. coli K12 corrected using 50x Illumina


SMRT-Assembly Results



Organism	Technology	Reference bp	Assembly bp	90% coverage	95% coverage/length	99%
<i>Leishmania major</i>	Illumina 150x 200bp	48,440	48,440	1	48,440 x 48,440	48,440 x 48,440 (200bp) 1
	PacBio-PbA 10x		48,440	1	48,440 x 48,440	48,440 x 48,440 (200bp) 1
	Both PacBio-PbA + Illumina 150x 200bp		48,440	1	48,440 x 48,440	48,440 x 48,440 (200bp) 1
<i>E. coli</i> K12	Illumina 150x 200bp	4,639,815	4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	PacBio-PbA 10x		4,639,815	10	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	Both PacBio-PbA + Illumina 150x 200bp		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
<i>E. coli</i> O157:H7	PacBio-PbA 10x	4,639,815	4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	PacBio-PbA 10x connected by 20x CCS		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	Both PacBio-PbA 10x + 20x CCS		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	PacBio-PbA 10x connected by 40x CCS		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	Both PacBio-PbA 10x + 40x CCS		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	Minimap2 (connected 40x CCS) assembly		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
<i>S. aureus</i> 8325a	Illumina 150x 200bp	11,111,111	11,111,111	95	11,111,111 x 11,111,111	11,111,111 x 11,111,111 (200bp) 1
	PacBio-PbA 10x		11,111,111	10	11,111,111 x 11,111,111	11,111,111 x 11,111,111 (200bp) 1
	Both PacBio-PbA 10x + Illumina 150x 200bp		11,111,111	95	11,111,111 x 11,111,111	11,111,111 x 11,111,111 (200bp) 1
<i>Mycobacterium tuberculosis</i>	Illumina 150x 200bp paired end (PE) 150x 200bp	4,639,815	4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	40x PE + 40x PE + 40x PE (paired end)		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1
	40x PE + 40x PE + 40x PE (paired end)		4,639,815	95	4,639,815 x 4,639,815	4,639,815 x 4,639,815 (200bp) 1

Hybrid assembly results using error corrected PacBio reads
Meets or beats Illumina-only or 454-only assembly in every case

Improved Gene Reconstruction




FOXP2 assembled on a single contig

Transcript Alignment



- Long-read single-molecule sequencing has potential to directly sequence full length transcripts
 - Raw reads and raw alignments (red) have many spurious indels inducing false frameshifts and other artifacts
 - Error corrected reads almost perfectly match the genome, pinpointing splice sites, identifying alternative splicing
- New collaboration with Gingeras Lab looking at splicing in human

Assembly Summary



Assembly quality depends on

1. **Coverage**: low coverage is mathematically hopeless
 2. **Repeat composition**: high repeat content is challenging
 3. **Read length**: longer reads help resolve repeats
 4. **Error rate**: errors reduce coverage, obscure true overlaps
- Assembly is a hierarchical, starting from individual reads, build high confidence contigs/unitigs, incorporate the mates to build scaffolds
 - Extensive error correction is the key to getting the best assembly possible from a given data set
 - Watch out for collapsed repeats & other misassemblies
 - Globally/Locally reassemble data from scratch with better parameters & stitch the 2 assemblies together




Whole Genome Alignment with MUMmer

Slides Courtesy of Adam M. Phillippy

amp@umics.umd.edu


Goal of WGA

- For two genomes, A and B , find a mapping from each position in A to its corresponding position in B



Not so fast...

- Genome *A* may have insertions, deletions, translocations, inversions, duplications or SNPs with respect to *B* (sometimes all of the above)




WGA visualization

- How can we visualize *whole* genome alignments?



- With an alignment dot plot

- $N \times M$ matrix

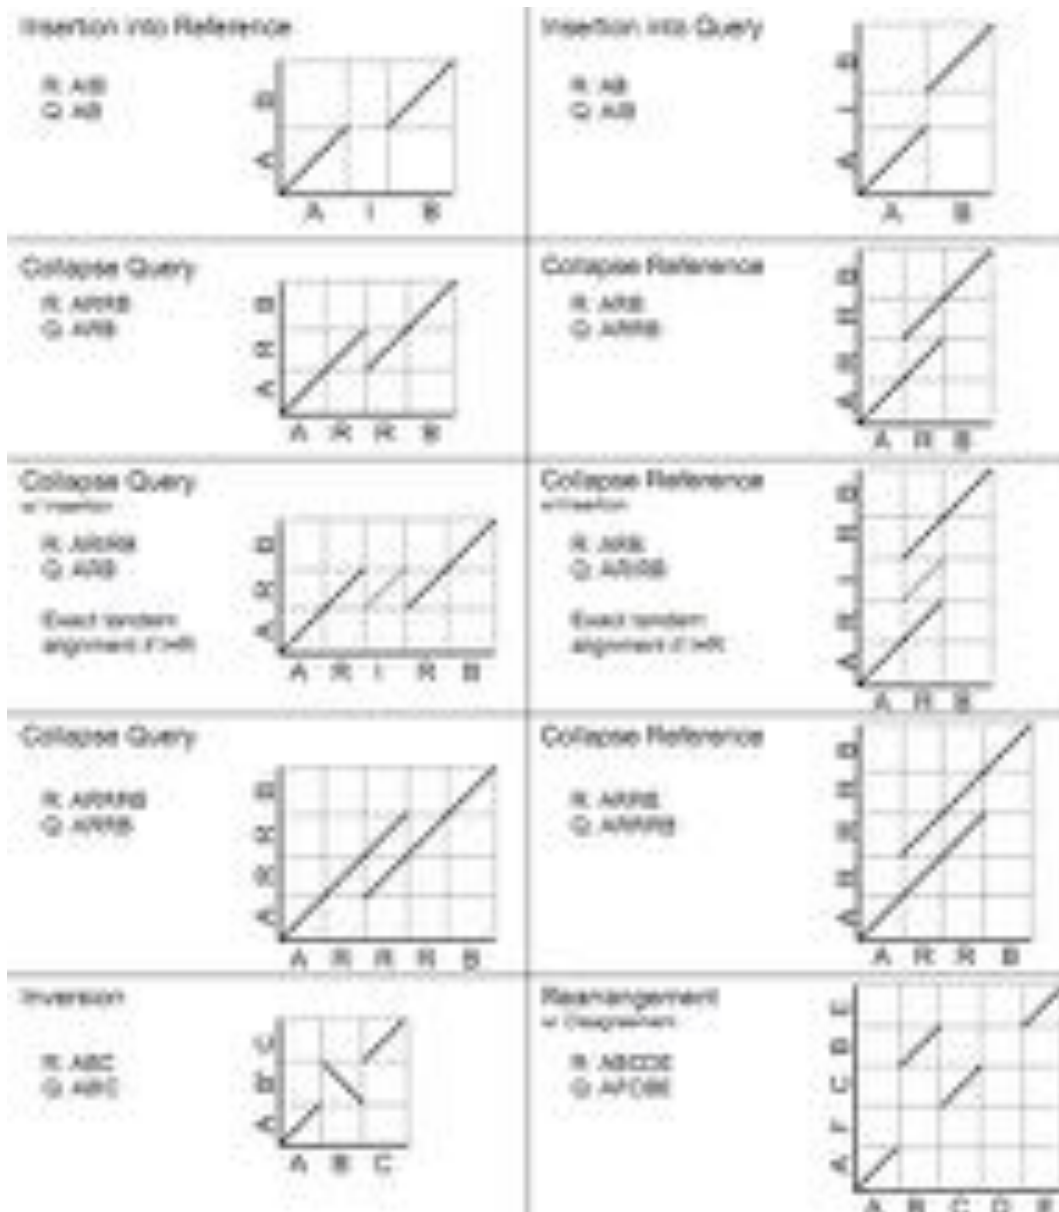
- Let i = position in genome A
 - Let j = position in genome B
 - Fill cell (i,j) if A_i shows similarity to B_j



- A perfect alignment between A and B would completely fill the positive diagonal



SV Types



- Different structural variation types / misassemblies will be apparent by their pattern of breakpoints
- Most breakpoints will be at or near repeats
- Things quickly get complicated in real genomes

<http://mummer.sf.net/manual/AlignmentTypes.pdf>

Seed-and-extend with MUMmer

How can quickly align two genomes?

1. Find maximal-unique-matches (MUMs)

- ♦ Match: exact match of a minimum length
- ♦ Maximal: cannot be extended in either direction without a mismatch
- ♦ Unique
 - ♦ occurs only once in both sequences (MUM)
 - ♦ occurs only once in a single sequence (MAM)
 - ♦ occurs one or more times in either sequence (MEM)

2. Cluster MUMs

- ♦ using size, gap and distance parameters

3. Extend clusters


- ♦ using modified Smith-Waterman algorithm

Fee Fi Fo Fum, is it a MAM, MEM or MUM?

MUM : maximal unique match _____

MAM : maximal almost-unique match - - - - -

MEM : maximal exact match




Seed and Extend

visualization

FIND all MUMs

CLUSTER consistent MUMs

EXTEND alignments



WGA example with **nucmer**

- *Yersina pestis* C092 vs. *Yersina pestis* KIM
 - High nucleotide similarity, 99.86%, but extensive reshuffling
 - High repeat content


```
nucmer -maxmatch C092.fasta KIM.fasta  
-maxmatch    Find maximal exact matches (MEMs)
```

```
delta-filter -m out.delta > out.filter.m  
-m    Many-to-many mapping
```

```
show-coords -r out.delta.m > out.coords  
-r    Sort alignments by reference position
```

```
dnadiff out.delta.m  
Construct catalog of sequence variations
```

```
mummerplot --large --layout out.delta.m  
--large      Large plot  
--layout    Nice layout for multi-fasta files
```



Review

Sequencing

1. Name 3 biological questions that can be answered using sequencing
2. Describe the overall process for identifying mutations in a genome using sequencing
 - Identifying de novo mutations
 - Measuring gene expression***
3. Suppose it takes 1000 hours to match 100M reads using the brute force algorithm against the human genome (3GB), how long would it take to search the barley genome (~6GB)?
 - wheat genome (~18GB), or pine tree genome (~24GB)?
 - Suppose it takes 10 hours using binary search against human, how long would it take for barley, wheat, or the pine tree?

Alignment

1. How many times do we expected GATTACA or GATTACA*2 or GATTACA*3 to be in the human genome?
 1. In the barley, wheat or pine tree genomes?
2. What is the suffix array for HURRICANESANDY
 1. Describe how I would find all occurrences of SAND in that suffix array
3. Describe how to find all occurrences of GATTACA in the human genome allowing at most 1 mismatch
4. What role do de novo mutations play in autism?

Assembly

1. Describe the overall process of genome assembly
2. What are the necessary data characteristics for a good genome assembly, and explain why they are necessary
3. Draw the de Bruijn graph using $k=1$ of the reads AR, BR, CR, RB, RC, RD and count the number of Eulerian paths
4. Draw the dot plot of GATTACA against GATTTTACA

Thank You!

<http://schatzlab.cshl.edu/>

